

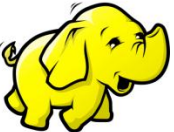
Real-Time Analytical Processing (RTAP) Using Spark and Shark

Jason Dai

Engineering Director & Principal Engineer
Intel Software and Services Group

Why Real-Time Analytical Processing (RTAP)?

Big data in large web sites

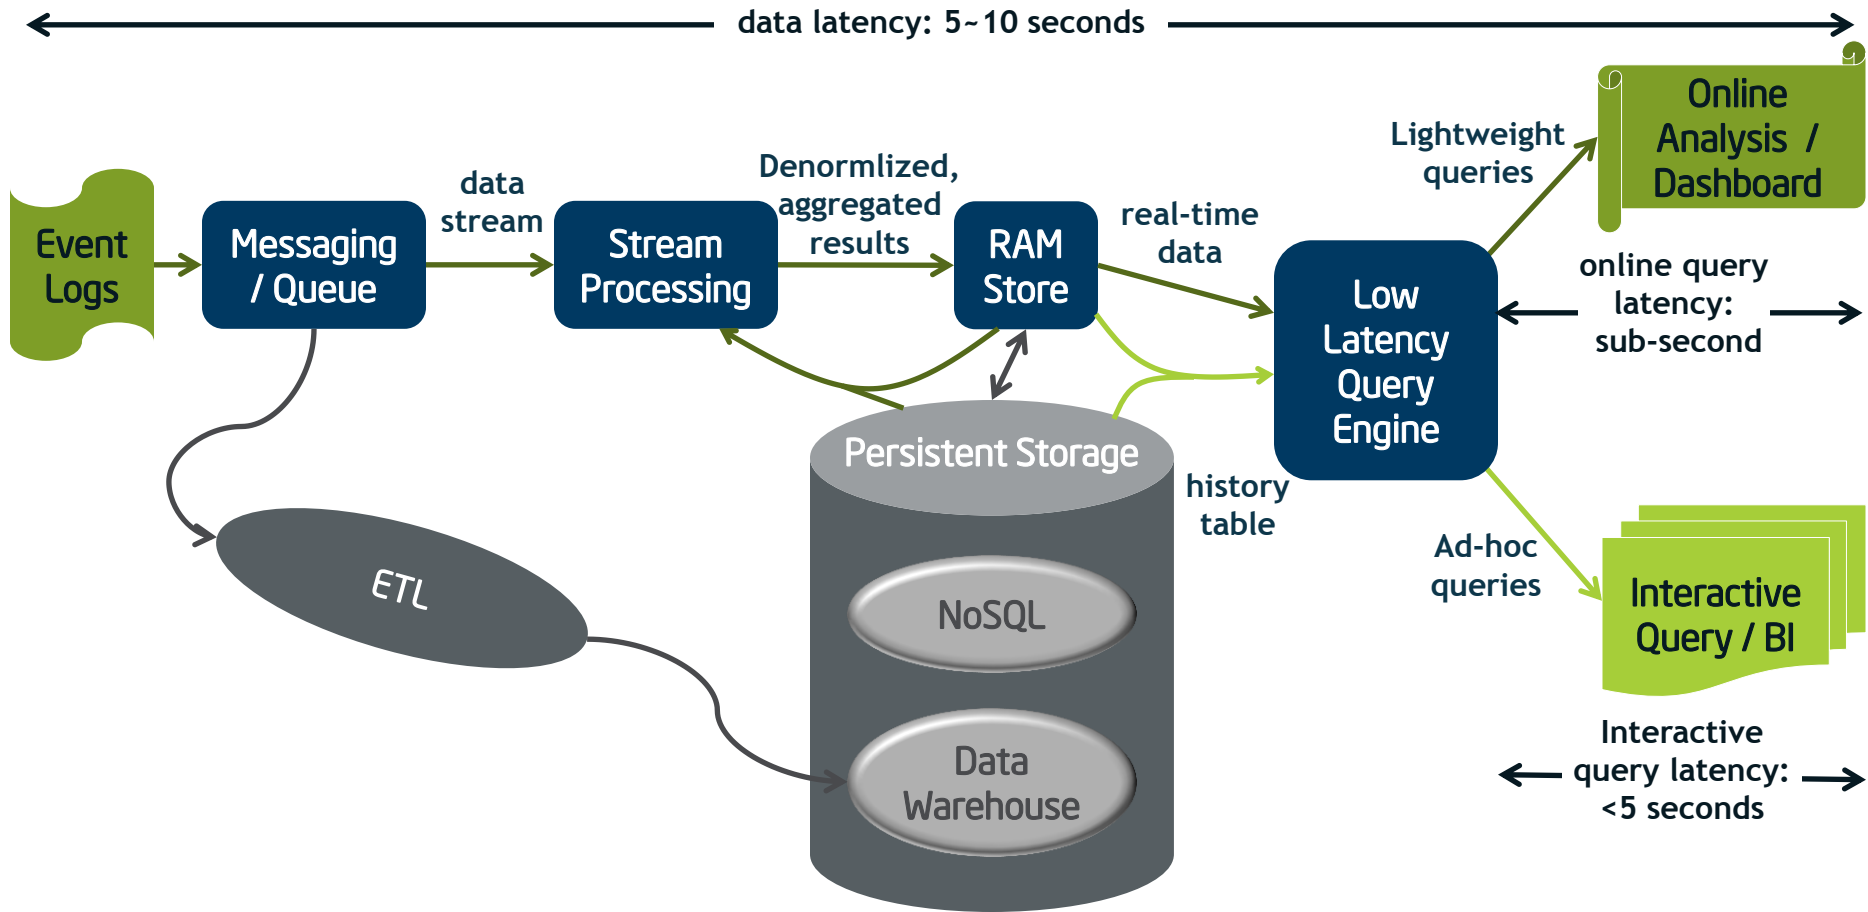
- The  is in the room

Moving beyond the elephant

- Discover and explore data iteratively and interactively for real-time insights
- *RTAP: Real-Time Analytical Processing*
 - Data continuously streamed in & processed in near real-time
 - Real-time data queried and presented in an online fashion
 - Real-time and history data combined and mined interactively
 - Predominantly RAM-based processing

We are partnering with several web sites in China
on building the *RTAP* framework using Spark & Shark


RTAP Architecture



RTAP Use Cases

Online dashboard

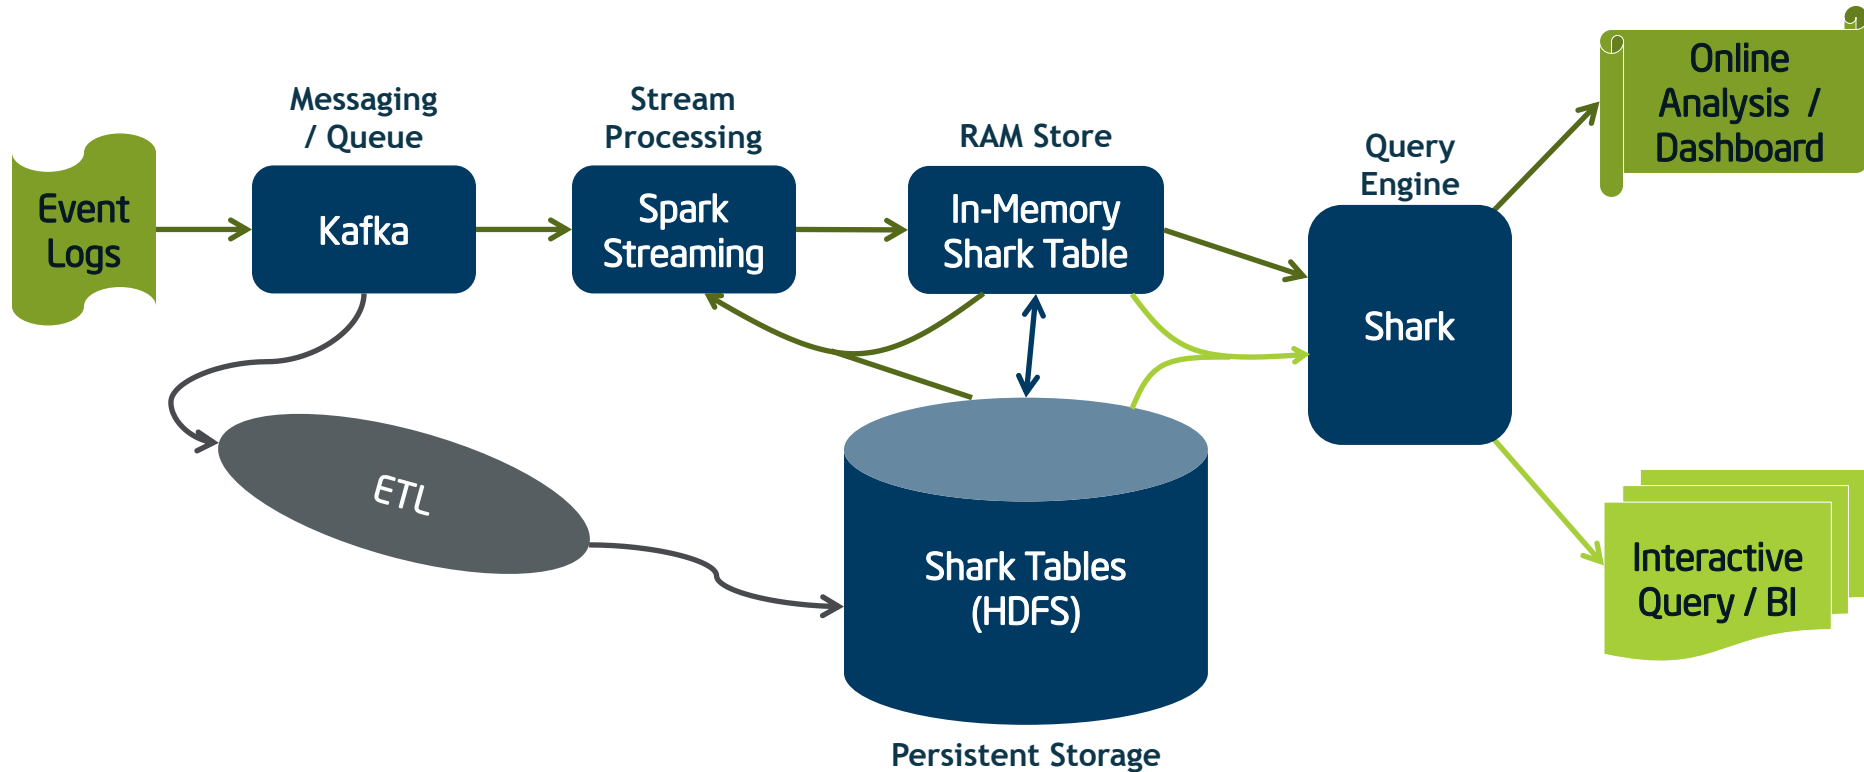
- Pages/Ads/Videos/Items – time base aggregations – break-down by categories/demography



Interactive BI

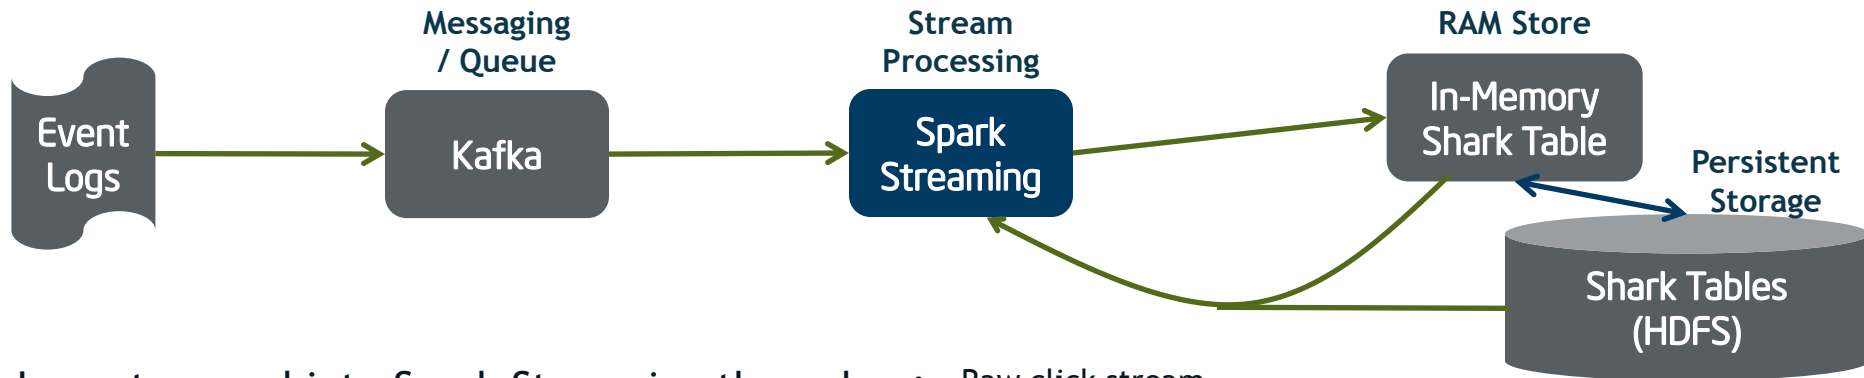
- Combined with history & dimension data when necessary
 - E.g., top 100 viewed videos under each category in the last month

RTAP Framework using Spark & Shark



A work in progress

Real-Time Data Stream Processing



Logs streamed into Spark Streaming through Kafka in real-time

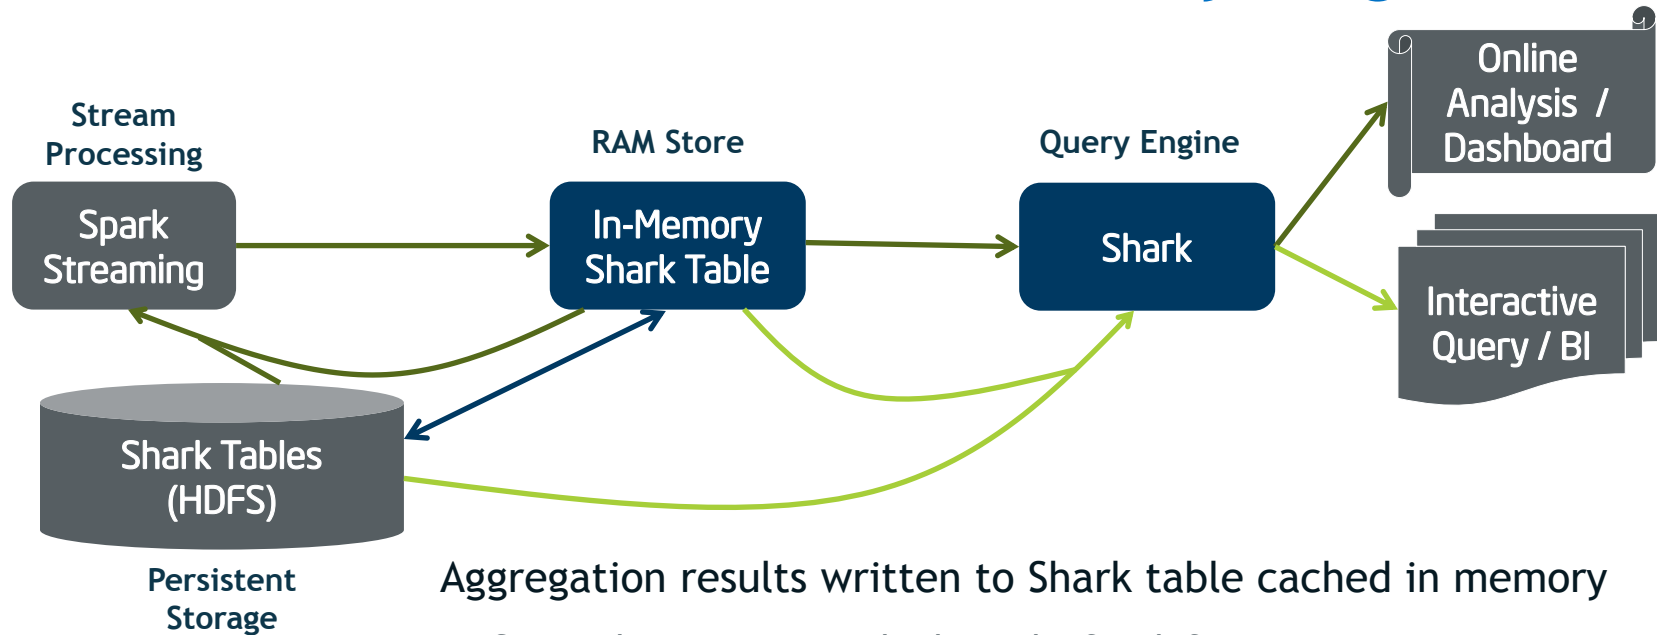
- Raw click stream
 - 0.6.38.68 - - BAF42487E0C7076CE576FAAB0E1852EC [14/Dec/2012 8:21:16 -0] "GET ?video=8745 HTTP/1.1" 101 1345 <http://www.foo.com/bar/?ivideo=8745> "Mozilla/4.0 (compatible; MSIE 5.5; Windows 98; Win 9x 4.90)"

Incoming logs processed by Spark Streaming in small batches (e.g., 10 seconds)

- Compute multiple aggregations over logs received in the last window (e.g., 1 minute)
- Join logs and history tables when necessary
- Compute page view in the last minute
 - E.g., www.foo.com/bar/?video=8745, www.foo.com/bar/, www.foo.com/, etc.
- Compute category view count in the last minute
 - E.g., join logs and the video table (assuming [video 8745](#) belongs to [/vehicle/car/sports](#)) for [/vehicle](#), [/vehicle/car](#), [/vehicle/car/sports](#), etc.

Plan to add the streaming support directly in Shark

Real-Time Data Store and Query Engine



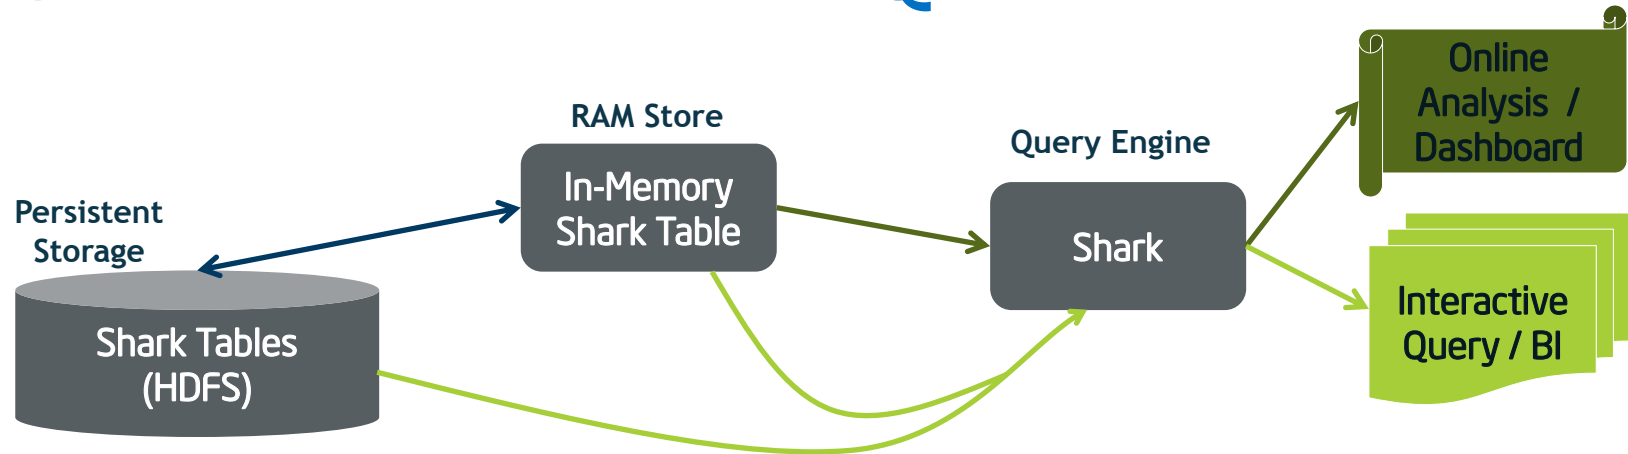
Aggregation results written to Shark table cached in memory

- Currently output as cached RDD by Spark Streaming
 - Require Spark Streaming embedded in the Shark server JVM
- Plan to move to Tachyon for better sharing and fault tolerance

Both real-time aggregations and history data queried through Shark

- History data loaded into memory for iterative mining
- Working on query optimizations & standard SQL-92 support

Online and Interactive Queries



Online analysis

- A lightweight UI frontending Shark for online dashboard
- Mostly time-based lightweight queries (filtering, ordering, TopN, aggregations, etc.) with sub-second latency

Interactive query / BI

- Ad-hoc, (more) complex SQL queries (with <5second latency)
- Heavily denormalized to eliminate join as much as possible

Experience and Current Work

Spark/Shark rocks!

- Lightweight, low latency & RAM-oriented
- Working on reliability & performance improvements
 - HA, isolation, fault tolerance, metadata handling, shuffle, etc.

Current improvements

- Shuffle performance improvements
 - Added a new netty-based shuffle module for Spark
- Shark query optimizations
 - Generated bytecode for expression evaluations
- Better operability
 - Added a new FairScheduler for Spark tasks
 - Added job history logs for Spark jobs

Future Work

- Integrate with Tachyon
- More shuffle improvements
 - Reduce the number of many (small) shuffle files
 - Actively remove shuffle files
 - Pipeline data shuffles and “map” tasks
- Further Shark query optimizations
 - Bytecode generation coverage (esp. aggregation)
 - Co-partitioned table for join
 - TopN pushdown

Future Work

- More Shark features
 - SQL-92 support (integrating “Project Panthera” with Shark)
 - Streaming support in Shark
- Even better operability
 - High availability

